

Chetan Aditya Lakka

AI Engineer | Agentic AI | LLM Evaluation & Reliability

✉ chetan.lakka@gmail.com ☎ +91 7659095665 🌐 chetanaditya.netlify.app 📄 Chetan Aditya 📱 ChetanAditya765

Profile

AI Engineer specializing in agentic AI, LLM evaluation and AI reliability with hands-on experience evaluating Computer Use Agents through RLHF/SFT-style workflows and structured QA. Skilled in analyzing multi-step agent behavior, hallucinations, grounding failures, tool and UI interactions, policy compliance, and model reliability. Built production-style AI systems using Python, FastAPI, LangGraph, PostgreSQL, pgvector, RAG, human-in-the-loop approvals, and automated evaluation with experience in multimodal AI and explainable ML.

Professional Experience

Turing

AI Quality Assurance Specialist - Gemini CUA

May 2026 – Present | Hyderabad (Remote)

- Evaluate browser-based Computer Use Agents across complex multi-step workflows, validating reasoning, UI interactions, tool execution, factual accuracy, task completion, and policy compliance.
- Analyze agent trajectories and recurring failure patterns across grounding, planning, execution, and recovery, translating observed failures into structured QA findings and reliability improvements.
- Design and refine evaluation rubrics, QA criteria, prompts, and regression workflows to improve agent reliability and identify behavioral regressions across evaluation cycles.
- Support evaluation automation and systematic failure analysis, strengthening the consistency and scalability of agent-quality workflows.

LLM Evaluation Specialist - Gemini CUA

Jan 2026 – May 2026 | Hyderabad (Remote)

- Performed RLHF/SFT-style side-by-side evaluations of browser-based AI agents using structured criteria for reasoning quality, instruction adherence, factual accuracy, and task completion.
- Evaluated multi-step agent trajectories to identify hallucinations, grounding failures, reasoning inconsistencies, and incorrect UI or tool interactions.
- Applied structured evaluation rubrics and documented recurring model failure modes to support model assessment, prompt refinement, and reliability analysis.

AI Training Data Specialist - OpenAI Video Analysis

Sep 2025 – Nov 2025 | Hyderabad (Remote)

- Worked on multimodal AI training data involving robotics tasks, human actions, and interactive video environments.
- Annotated timestamped action and interaction sequences supporting temporal grounding and spatial reasoning in vision-language systems.
- Evaluated multimodal model outputs to identify reasoning gaps, inconsistencies, and failure cases.
- Performed structured quality checks supporting training-data quality and evaluation reliability.

Projects

ResolveAI - Agentic AI Reliability & IT Access Automation 📄, Python | FastAPI | LangGraph | PostgreSQL | pgvector | RAG | Next.js | Docker

- Built a production-style agentic AI system that converts repository-access requests into policy-grounded workflows with durable execution, human approval, typed tool actions, independent verification, and automated evaluation.
- Designed a persisted LangGraph workflow combining RAG-based policy retrieval, deterministic authorization, resumable human-in-the-loop approvals, safe escalation, and auditable tool execution.
- Implemented reliability controls including idempotent operations, least-privilege rules, action-boundary revalidation, independent read-after-write verification, and failure-safe ticket resolution.
- Developed an independent evaluation framework covering policy evidence, resource grounding, tool scope, approval ordering, verification, escalation, and terminal-state correctness.
- Validated the local release with 482 backend tests, 12 Chromium E2E tests, and 37/37 evaluation/deployment scenarios, including persistence of workflow state across container restarts.

Reverse Attribution - ML Failure Analysis & Explainability 📄, Python | PyTorch | Captum | BERT | RoBERTa | ResNet | XAI

- Developed Reverse Attribution, a supervised post-hoc failure-analysis framework that identifies counter-evidence by analyzing features with negative attribution toward the known true class.
- Designed Counterfactual Pull and Attribution-Flip (A-Flip) metrics to measure true-class logit recovery under feature interventions and disagreement between competing class explanations.
- Built a reproducible frozen-checkpoint evaluation pipeline across NLP and vision models, generating 232 corrected per-example analyses with deterministic sampling, provenance tracking, visualization, and 123 passing tests.
- Audited and repaired the research implementation to align equations, attribution logic, experimental outputs, figures, and manuscript claims while preserving historical model checkpoints.

CereBro - AI-Enabled Cognitive Screening Application 📄, Android | Python | LangChain | REST APIs | On-Device ML | Cloud Integration

- Built an AI-enabled Android application supporting screening workflows for autism, dyslexia, and dementia using on-device machine learning.
- Integrated LLM-based analysis using Python and LangChain for automated response interpretation and feedback generation.
- Designed REST API and cloud integration workflows connecting mobile inference, AI services, and application logic.

Skills

- **AI / LLM Systems** — Agentic AI, Large Language Models, LangGraph, LangChain, RAG, LLM Agents, Prompt Engineering, OpenAI API, LLM Evaluation, RLHF, SFT, AI Reliability, Multimodal AI, Hugging Face
- **Programming / ML** — Python, PyTorch, TensorFlow, Scikit-learn, NumPy, Pandas, OpenCV
- **Backend / Data** — FastAPI, REST APIs, PostgreSQL, pgvector, SQLAlchemy, Pydantic, Alembic, SQL, JSON
- **Testing / Engineering** — Docker, Docker Compose, Pytest, Playwright, Git, API Integration, Model Deployment
- **Explainability / Evaluation** — SHAP, Integrated Gradients, Captum, Model Interpretability, Failure Analysis, Evaluation Design

Education

B.Tech in Electronics & Computer Engineering, Sreenidhi Institute Of Science And Technology

2021 – 2025

Selected Publications

Reverse Attribution: Explaining Model Failures via Counter-Evidence, NCNTAIA Conference

CereBro: A Cloud-Integrated Android Application for Cognitive Disorder Screening, NCNTAIA Conference

Certifications

Artificial Intelligence: Knowledge Representation And Reasoning [🔗](#), NPTEL

LLM Trainer Certification [🔗](#), micro1

AWS Academy Machine Learning Foundations [🔗](#), Amazon Web Services (AWS)

AWS Academy Cloud Foundations [🔗](#), Amazon Web Services (AWS)